




Large

$S = \text{Where are we going}$

Language

$$P(S) = P(\text{Where}) \times P(\text{are} \mid \text{Where}) \times P(\text{we} \mid \text{Where are}) \times P(\text{going} \mid \text{Where are we})$$



Models

CIS 7000 - Fall 2024

Introduction

Professor Mayur Naik



Today's Agenda

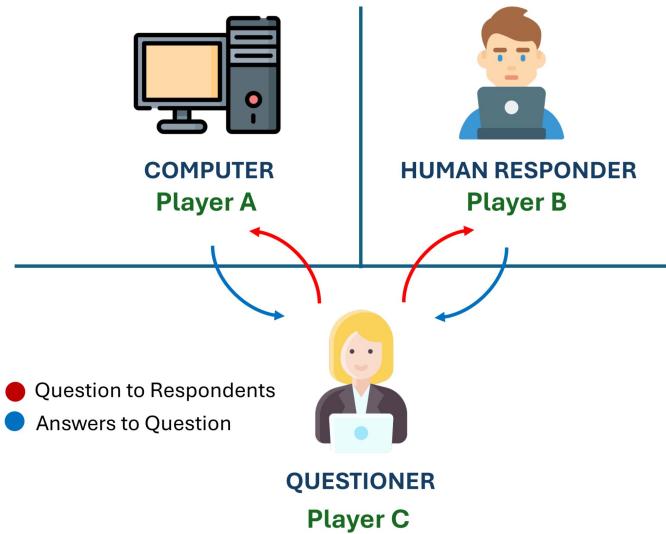
- The Turing Test
- Overview of LLMs
 - How do LLMs work, What LLMs can do, Limitations of LLMs, What is the future
- Course Logistics

The Imitation Game (aka The Turing Test)

Proposed in 1950 by Alan M. Turing who is considered the father of theoretical computer science.

Tests a machine's ability to exhibit intelligent behaviour equivalent to, or indistinguishable from, that of a human – via language.

Language modeling has since been proposed as a benchmark to measure progress toward AI.



"I believe that in about fifty years' time it will be possible to programme computers, with a storage capacity of about 10^9 , to make them play the imitation game so well that an average interrogator will not have more than 70% chance of making the right identification after five minutes of questioning."

– A. Turing. Computing Machinery and Intelligence. Mind, 1950.

Eras of Language Modeling

Symbolic Era

Pre-1990

Rule-based approaches

Expert systems

Limited generalization

Statistical Era

1990-2010

Data-driven approaches

Probabilistic models

Introduction of corpora

Scale Era

2010 onwards

Deep learning and neural nets

General-purpose LMs

Massive datasets and compute

Turing Test

ELIZA

ChatGPT

1950

1966

2022



ELIZA (1966)

Welcome to

EEEEEE	LL	IIII	ZZZZZZ	AAAAAA
EE	LL	II	ZZ	AA AA
EEEEEE	LL	II	ZZZ	AAAAAAA
EE	LL	II	ZZ	AA AA
EEEEEE	LLLLL	IIII	ZZZZZZ	AA AA

Eliza is a mock Rogerian psychotherapist.

The original program was described by Joseph Weizenbaum in 1966.
This implementation by Norbert Landsteiner 2005.

ELIZA: Is something troubling you ?

YOU: Men are all alike.

ELIZA: What is the connection, do you suppose ?

YOU: They're always bugging us about something or other.

ELIZA: Can you think of a specific example ?

YOU: Well, my boyfriend made me come here.

ELIZA: Is it important to you that your boyfriend made you come here ?

YOU: He says I'm depressed much of the time.

ELIZA: I am sorry to hear that you are depressed.

YOU: It's true. I am unhappy.

ELIZA: Can you explain what made you unhappy ?

YOU:

Early NLP program developed by Joseph Weizenbaum at MIT.

Created illusion of a conversation by rephrasing user statements as questions using pattern matching and substitution methodology.

One of the first programs capable of attempting the Turing test.

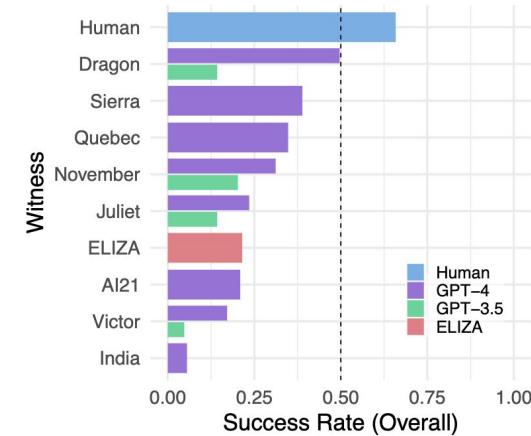

Try it out at <https://web.njit.edu/~ronkowitz/eliza.html>

Has AI Passed The Turing Test?

How do we even tell?

"The best-performing GPT-4 prompt passed in 49.7% of games, outperforming ELIZA (22%) and GPT-3.5 (20%), but falling short of the baseline set by human participants (66%)."

C. Jones and B. Bergen. [Does GPT-4 pass the Turing test?](#) 2024.



"ChatGPT-4 exhibits behavioral and personality traits that are statistically indistinguishable from a random human from tens of thousands of human subjects from more than 50 countries."

Q. Mei et al. [A Turing test of whether AI chatbots are behaviorally similar to humans](#). PNAS, 2024.

A Social Turing Game

Chat with someone for two minutes and guess if it was a fellow human or an AI bot. The AI bots in the game are chosen from a mix of different LLMs, including Jurassic-2, GPT-4, Claude, and Cohere.

<https://www.humanornot.ai/>

Part of a larger scientific research project by AI21 Labs.

D. Jannai et al. [Human or Not? A Gamified Approach to the Turing Test](#). 2023.

Question: Can you identify a flaw of using this game as a Turing Test?

Has AI Passed The Turing Test?

How do we even tell?

Is the test even a valid measure of AI's capabilities?

What are the ethical implications of passing the test?

And many others ...

Overview of LLMs

How do LLMs work

What is the technology underlying a chatbot like chatGPT?

What LLMs can do

What functionality beyond chatbots does the technology enable?

Limitations of LLMs

What fundamental challenges remain to be addressed?

What is the Future

How is research addressing those challenges?

How Do LLMs Work

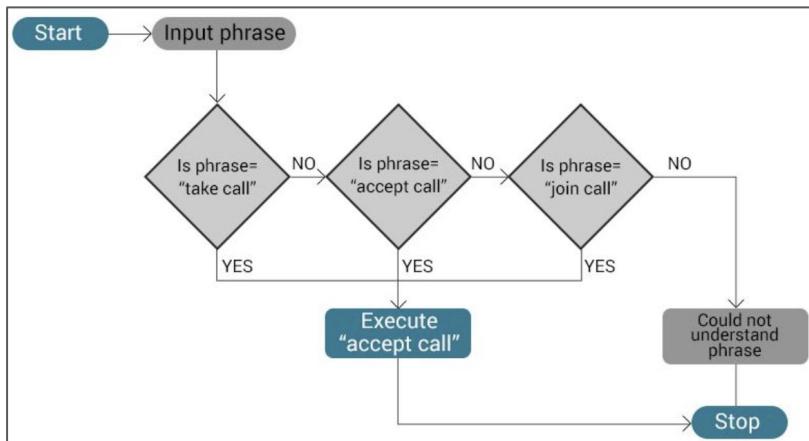
Let's Take a History Tour!

“Those who cannot remember the past are condemned to repeat it.”

— George Santayana. *The Life of Reason*, 1905.

Linguistic Foundations

Rule-based approaches



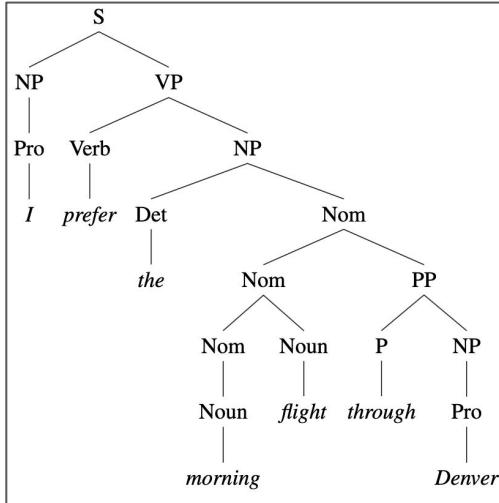
Example rule in a chatbot based on AIML (Artificial Intelligence Markup Language) which was developed in 1992-2002.

AIML formed the basis for a highly extended Eliza called A.L.I.C.E. ("Artificial Linguistic Internet Computer Entity").

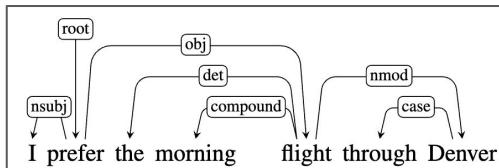
Linguistic Foundations

Semantic parsing: analyzing the linguistic structure of text

Example of
constituency parsing
using a
context-free grammar.



Same example using
dependency parsing.



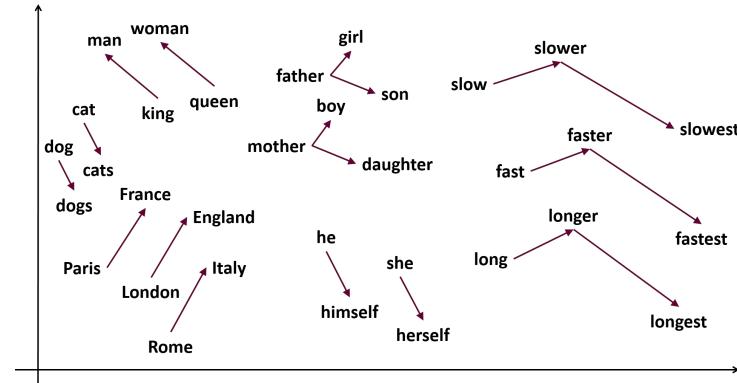
The introduction of corpora ...

The **Penn Treebank (PTB)** corpus developed during 1989-1996 was widely used for evaluating models for sequence labelling. The task consists of annotating each word with its Part-of-Speech tag.

M. Marcus et al. [Building a Large Annotated Corpus of English: The Penn Treebank](#). Computational Linguistics, 1993.

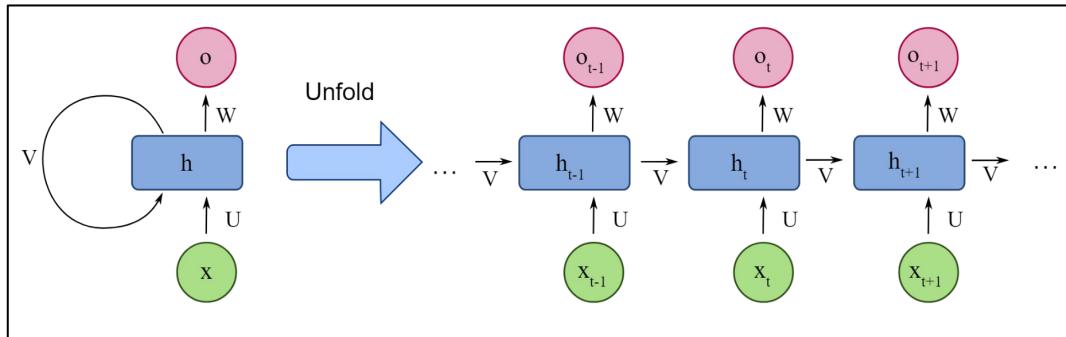
Word Embeddings

- Represent each word as a “vector” of numbers.
- Converts a “discrete” representation to “continuous”, allowing for:
 - More “fine-grained” representations of words.
 - Useful computations such as cosine and Euclidean distance.
 - Visualization and mapping of words onto a semantic space.
- Can be learnt in self-supervised manner from a large corpus.
- Examples:
 - Word2Vec (2013), GloVe, BERT, ELMo

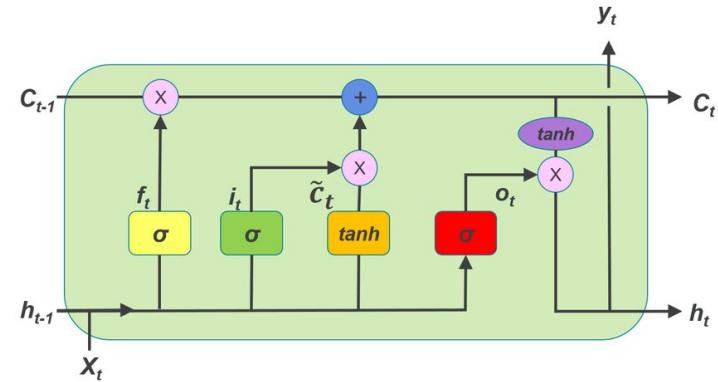


Seq2Seq Models

- Recurrent Neural Networks (RNNs)
- Long Short-Term Memory Networks (LSTMs)
- Capture dependencies between input tokens
- Gates control the flow of information



A simple RNN shown unrolled in time. Network layers are recalculated for each time step, while weights U, V and W are shared across all time steps.




A single LSTM unit displayed as a computation graph. The inputs to each unit consists of the current input x_t , previous hidden state h_{t-1} , and previous context c_{t-1} . The outputs are a new hidden state h_t and an updated context c_t .

Self-Attention and Transformers

- Allows to “focus attention” on particular aspects of the input while generating the output.
- Done by using a set of parameters, called “weights,” that determine how much attention should be paid to each input token at each time step.
- These weights are computed using a combination of the input and the current hidden state of the model.

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

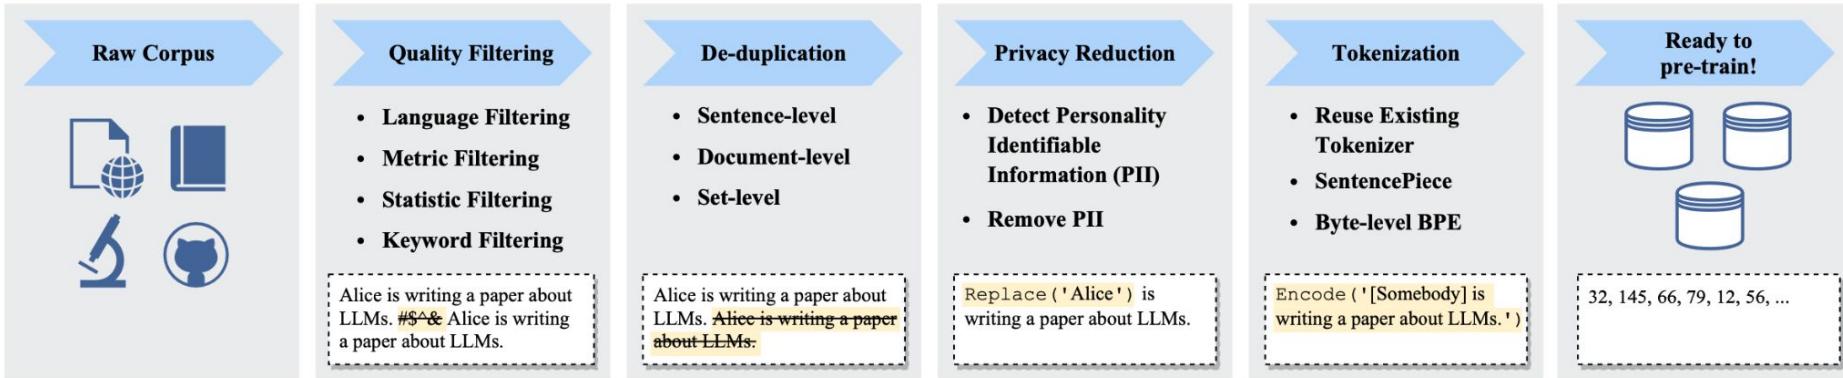
A. Vaswani et al. [Attention Is All You Need](#). NeurIPS 2017.



In encoding the word "it", one attention head is focusing most on "the animal", while another is focusing on "tired". The model's representation of the word "it" thus bakes in some of the representation of both "animal" and "tired".
<https://jalammar.github.io/illustrated-transformer/>

Pre-Training: Data Preparation

A typical data preparation pipeline for pre-training LLMs:



W. Zhao et al. [A Survey of Large Language Models](#). 2023.

Pre-Training Data Quality Reduces Reliance on Compute

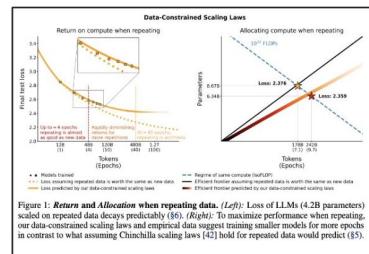
Recent work finds smaller amounts of higher quality data removes the need for a larger model.

There is increasing evidence that efforts to better curate training corpus, including **deduping, pruning data and investing in synthetic data** can compensate for the need for larger networks and/or improve training dynamics.

	% train examples with dup in train	% valid with dup in train	% valid in train
C4	3.04%	1.59%	4.60%
RealNews	13.63%	1.25%	14.35%
LM1B	4.86%	0.07%	4.92%
Wiki40B	0.39%	0.26%	0.72%

Table 2: The fraction of examples identified by NEARDUP as near-duplicates.

[Lee et al. 2022](#)



↳ Cohere For AI

[Muennighoff et al.
2023](#)

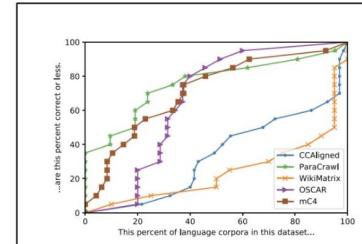



Figure 1: Fraction of languages in each dataset below a given quality threshold (percent correct).

[Kreutzer et al. 2022](#)

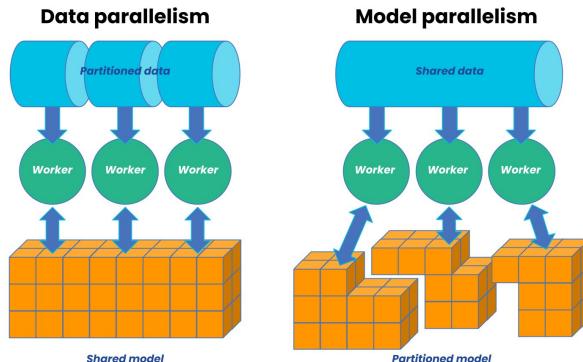


[Marion et al. 2023](#)


Pre-Training: Parallelism

4D Parallelism to minimize bottlenecks and maximizes efficiency: combines Data, Context, Pipeline (Vertical), and Tensor (Horizontal) Parallelism.

- **Data Parallelism** parallelizes tasks to speed up data processing and model iterations.
- **Context Parallelism** splits input sequences into chunks to be processed separately.



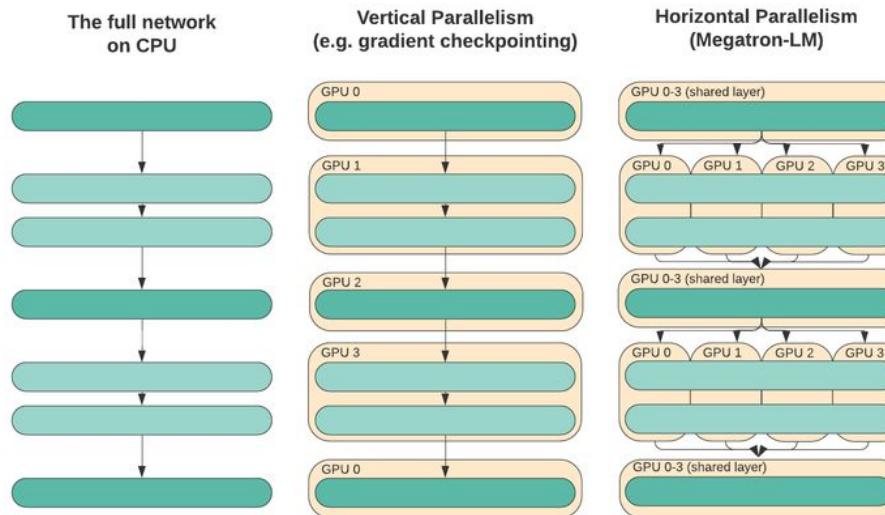
K. Pijanowski and M. Galarnyk.
[What is Distributed Training?](#) 2022.



S. Li et al. [Sequence Parallelism: Long Sequence Training from System Perspective](#). 2021.

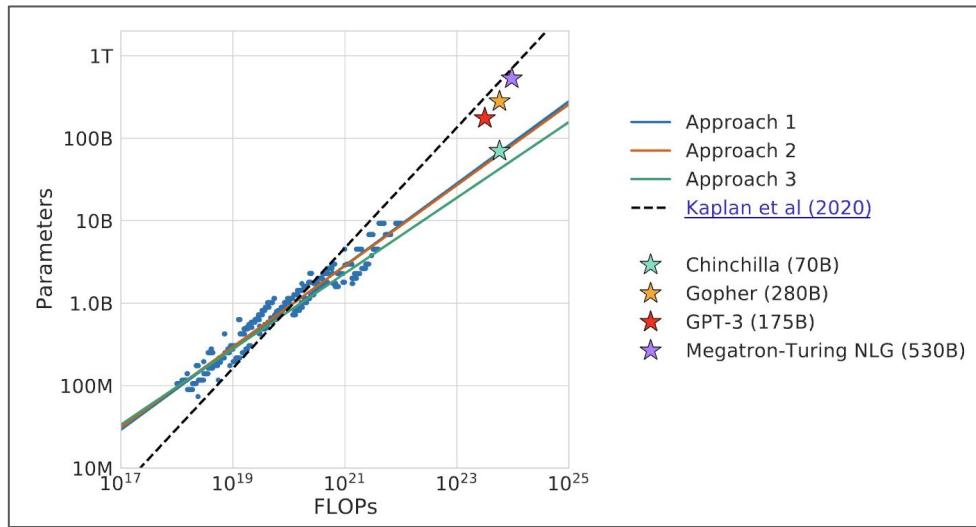
Pre-Training: Parallelism

- **Pipeline Parallelism** separates a model based on its layers, allowing higher throughput.
- **Tensor Parallelism** splits matrices across GPUs to reduce peak memory consumption.



Pre-Training: Scaling Laws

Given a fixed compute budget, what is the optimal model size and training dataset size for training a transformer LM?



Chinchilla Scaling Law:

For every doubling of model size, the number of training tokens must also be doubled.

Post-Training: Instruction-Tuning and Alignment

1

Pre-Training

2

Instruction Fine-tuning

3

Reinforcement Learning from Human Feedback

Massive amounts of data from Internet, books, etc.

Teach model to respond to instructions.

Problem: A model that can babble on about anything, but not aligned with what we want (e.g. Question-Answering)

<USER>

Can you write a short introduction about the relevance of the term "monopsony" in economics? Please use examples related to potential monopsonies in the labour market and cite relevant research.

<ASSISTANT>

"Monopsony" refers to a market structure where there is only one buyer for a particular good or service. In economics, this term is particularly relevant in the labor market, where a monopsony employer has significant power over the wages and working conditions of their employees. The presence of a monopsony can result in lower wages and reduced employment opportunities for workers, as the employer has little incentive to increase wages or provide better working conditions...

Teach model to produce output closer to what humans like.

You
Write a haiku about paperclips

ChatGPT
Steel curves intertwine,
Holding life's pages in line,
Quiet, small design.

ChatGPT
Silver arcs entwined,
Humble joiner of pages,
Quiet strength, aligned.

ChatGPT
Silver arcs entwined,
Holding life's pages, aligned,
In steel we find, bind.

Evaluation

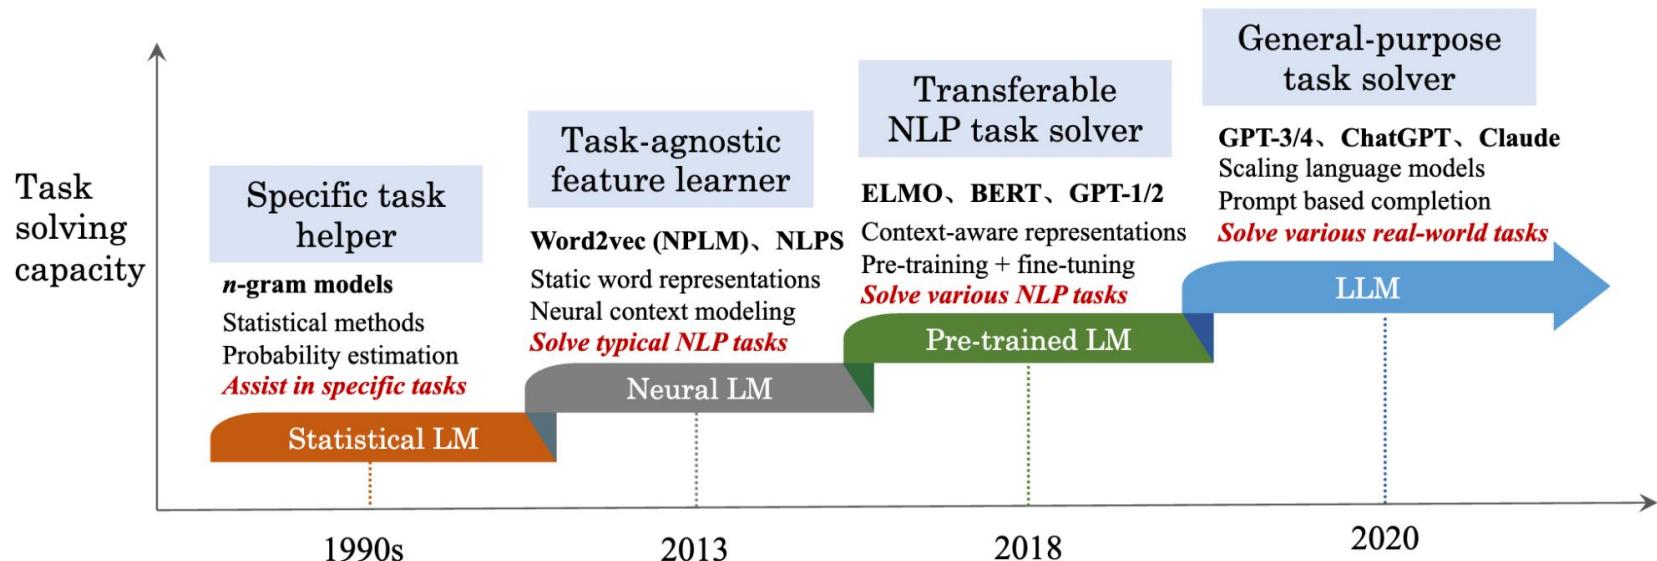
- Datasets
 - GLUE, SuperGLUE (General language understanding)
 - HumanEval (Coding)
 - HellaSwag (Commonsense reasoning)
 - GSM-8K (Math)
- Human Preferences
 - [Chatbot Arena](#): Crowdsourced platform where humans vote on pairwise comparisons of different LLMs (akin to Elo rating system in Chess).
- LLMs as Judges
 - LLM can approximate human preference with far lower cost!
 - L. Zheng et al. [Judging LLM-as-a-Judge with MT-Bench and Chatbot Arena](#). NeurIPS 2023 Datasets and Benchmarks Track.

How Do LLMs Work: Key Topics

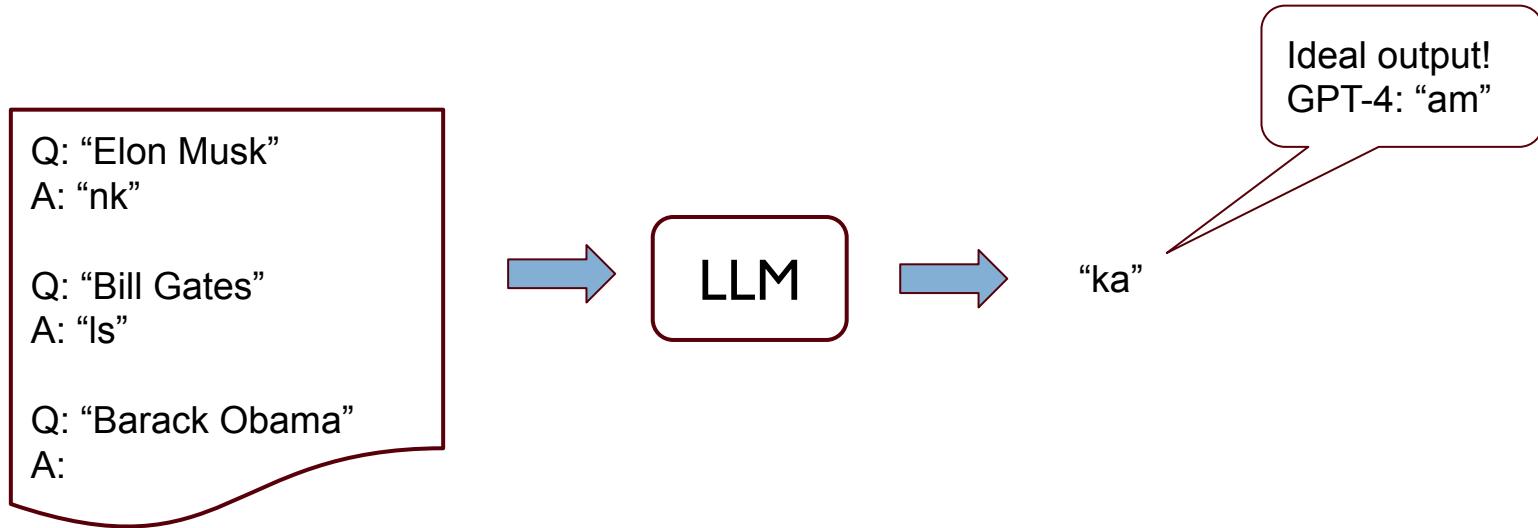
- Transformer Architecture
 - Self-Attention, Input/Output Processing, Architecture Variations, Training and Inference
- Pre-Training
 - Data Preparation (Tokenization, etc.), Parallelism, Scaling Laws
- Post-Training
 - Instruction Following/Tuning, Alignment
- Evaluation

What LLMs Can Do

Evolution of LMs from Perspective of Task-Solving Capacity

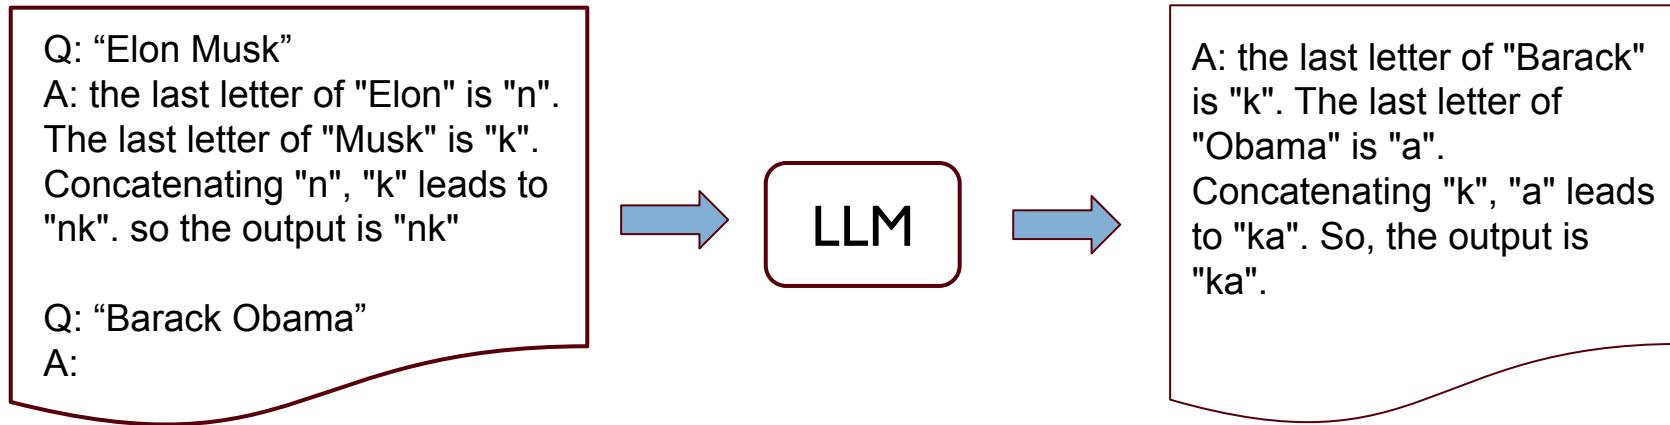


Few-Shot Prompting



T. Brown et al. [Language Models are Few-Shot Learners](#). NeurIPS 2020.

Chain-of-Thought Prompting



CoT as an Emergent Property of Model Scale

StrategyQA

Q: Yes or no: Would a pear sink in water?

A: The density of a pear is about 0.6 g/cm³, which is less than water. Thus, a pear would float. So the answer is no.

Date Understanding

Q: The concert was scheduled to be on 06/01/1943, but was delayed by one day to today. What is the date 10 days ago in MM/DD/YYYY?

A: One day after 06/01/1943 is 06/02/1943, so today is 06/02/1943. 10 days before today is 05/23/1943. So the answer is 05/23/1943.

SayCan (Instructing a robot)

Human: How would you bring me something that isn't a fruit?

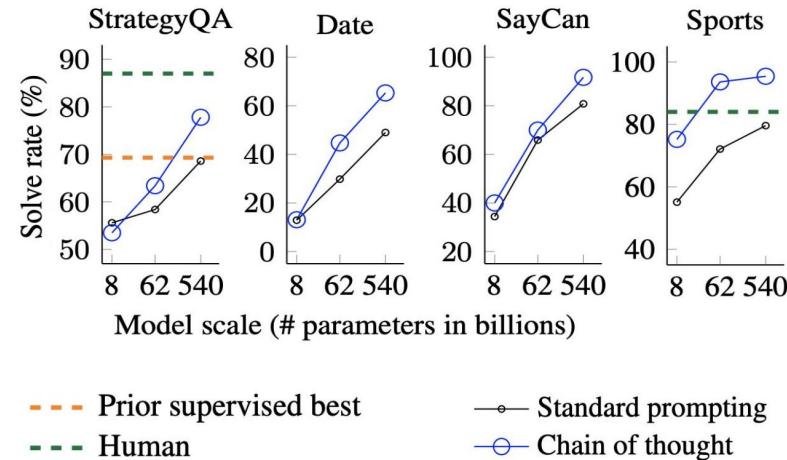
Explanation: the user wants something to eat that isn't a fruit. An energy bar is not a fruit, so I will bring the user an energy bar.

Plan: 1. find(energy bar) 2. pick(energy bar) 3. find(user) 4. put(energy bar) 5. done().

Sports Understanding

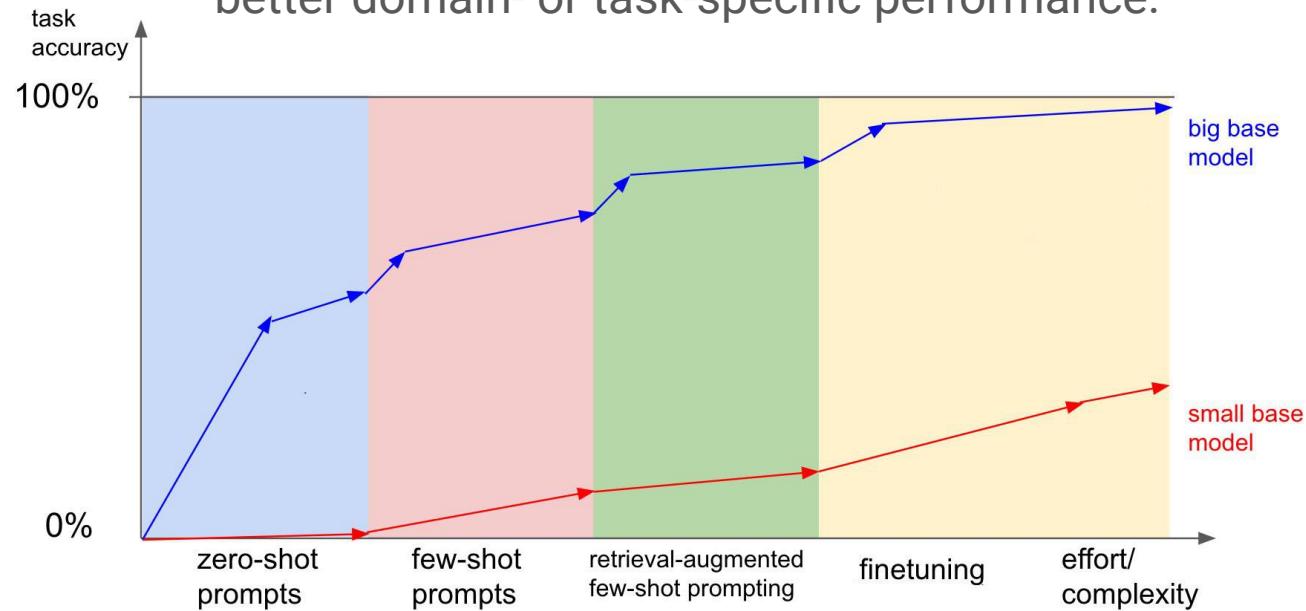
Q: Is the following sentence plausible? "Joao Moutinho caught the screen pass in the NFC championship."

A: Joao Moutinho is a soccer player. The NFC championship is part of American football, not soccer. So the answer is no.



From Prompting to Fine-Tuning

Unlike prompting, fine-tuning actually changes the model under the hood, giving better domain- or task-specific performance.



Source: Andrej Karpathy @karpathy (not to scale)

Case Study in Law: Harvey AI

- Startup building a custom-trained case law model for drafting documents, answering questions about complex litigation scenarios, and identifying material discrepancies between hundreds of contracts.
- Added **10 billion tokens** worth of data to power the model, starting with case law from Delaware, and then expanding to include all of U.S. case law.
- Attorneys from 10 large law firms preferred custom model's output versus GPT-4's **97% of the time**. Main benefit was reduced hallucinations!

[Open AI Customer Stories: Harvey](#). April 2024.

Case Study in Law: Harvey AI

Prompt: What is a claim of disloyalty?

GPT-4

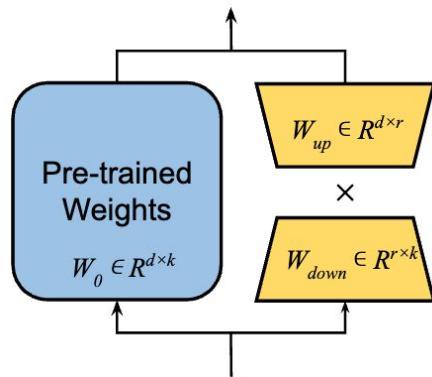
Harvey Custom Model

GPT-4

Custom-Trained Model

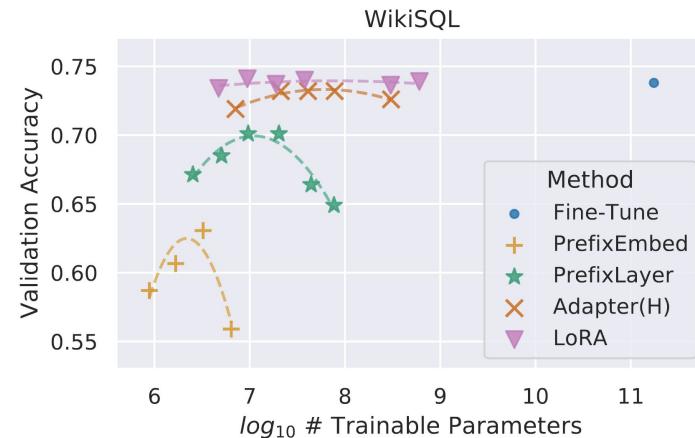
Preferred 97% of the time

Parameter Efficient Fine-Tuning (PEFT)

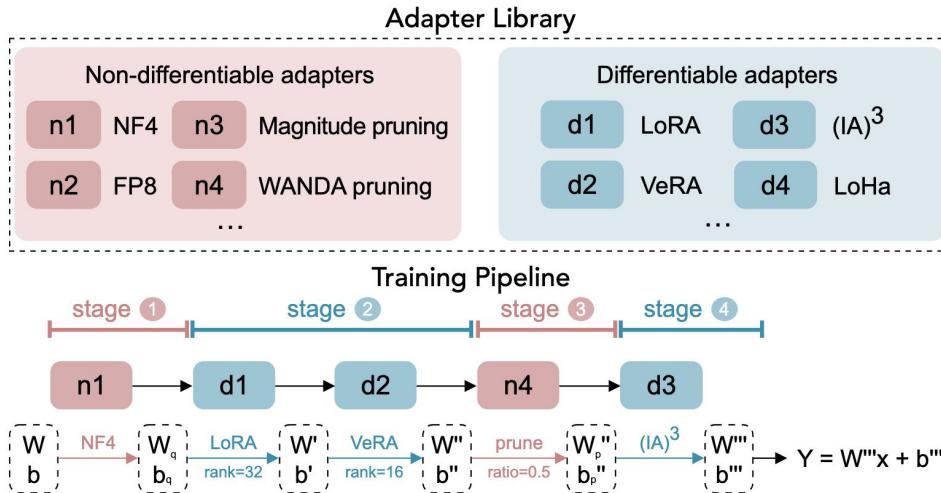


Techniques like LoRA construct a **low-rank** parameterization for parameter efficiency during training. For inference, the model can be converted to its original weight parameterization to ensure unchanged inference speed. Blue represents frozen while yellow represents trainable parameters.

GPT-3 175B validation accuracy vs. number of trainable parameters of several adaptation methods on WikiSQL. LoRA exhibits better scalability and task performance.



Design Spaces and The CLAM Framework



CLAM enables unlimited chaining of popular optimization techniques in parameter-efficient finetuning, quantization, and pruning on nearly every modern LLM.

N. Velingker et al. CLAM: Unifying Finetuning, Quantization, and Pruning through Unrestricted Chaining of LLM Adapter Modules. 2024.

What LLMs Can Do: Key Topics

- Prompt Engineering
 - Few-Shot, Chain-of-Thought (CoT), etc.
- Adaptation (aka Fine-Tuning)
 - Parameter-Efficient Techniques (PEFT)
 - Design Spaces
 - The CLAM Framework

Limitations of LLMs

Unreliable Reasoning Even On Simple Tasks

how many r's are there in the word
"strawberry"?



There are two "r's" in the word "strawberry."

Probably due to **tokenization!**

9.11 and 9.9, which one is larger?

9.11 and 9.9, which one is larger?

Generated by gpt-4o's tokenizer.




9.11 is larger than 9.9.

Try it out at:

<https://tiktoktokenizer.vercel.app/>

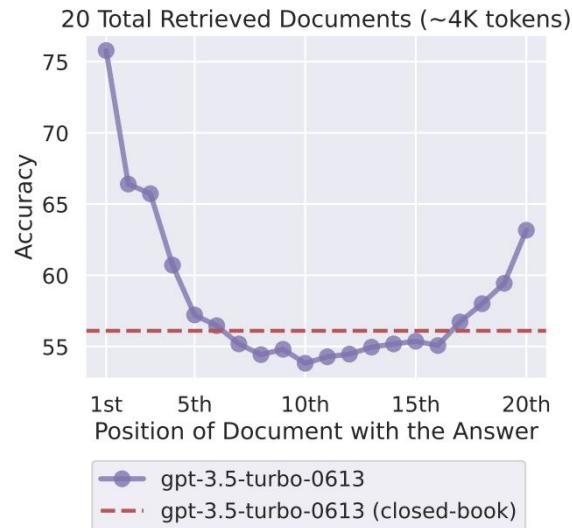
Jailbreaking Can Bypass Safety



Process of manipulating prompts to bypass an LLM's safeguards, leading to harmful outputs.

PAIR—which is inspired by social engineering attacks—uses an attacker LLM to automatically generate jailbreaks for a separate targeted LLM. The attacker LLM iteratively queries the target LLM to update and refine a candidate jailbreak, often in fewer than twenty queries.

Long Contexts Can Hurt Accuracy



Changing the location of relevant information within the model’s input context results in a U-shaped performance curve—models are better at using relevant information that occurs at the very beginning (primacy bias) or end of its input context (recency bias), and performance degrades significantly when models must access and use information located in the middle of its input context.

Limitations of LLMs: Key Topics

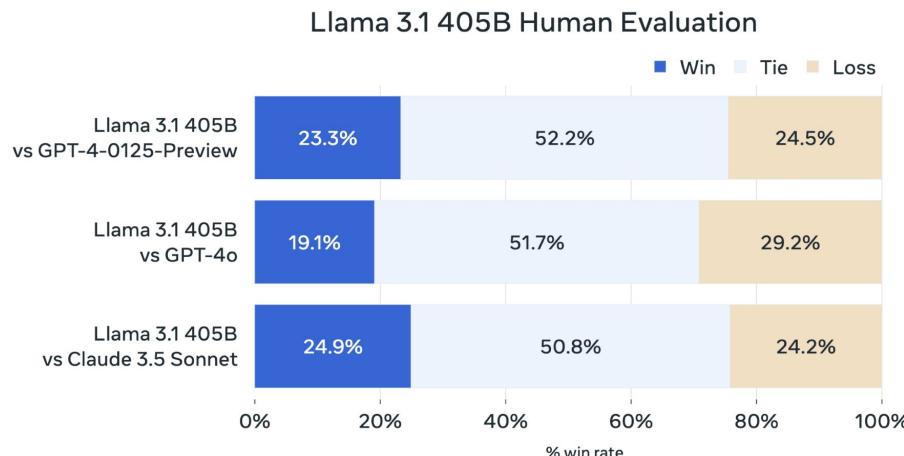
- Reasoning and Planning
- Hallucinations
- Limited Context
- Safety
- Interpretability
- Cost and Energy

What Is The Future

Synthetic Data / Distillation

*"In addition to having significantly better cost/performance relative to closed models, the fact that the 405B model is open will make it the best choice for **fine-tuning** and **distilling** smaller models."*

– M. Zuckerberg. [Open Source AI Is the Path Forward | Meta](#). 2024.





Rank* (UB)	Model	Arena Score
1	ChatGPT-4o-latest (2024-08-08)	1317
2	Gemini-1.5-Pro-Exp-0801	1298
2	Grok-2-08-13	1293
3	GPT-4o-2024-05-13	1286
5	GPT-4o-mini-2024-07-18	1275
5	Claude 3.5 Sonnet	1271
5	Grok-2-Mini-08-13	1268
6	Gemini Advanced App (2024-05-14)	1267
6	Meta-Llama-3.1-405b-Instruct	1266
7	GPT-4o-2024-08-06	1262

<https://larena.ai/>

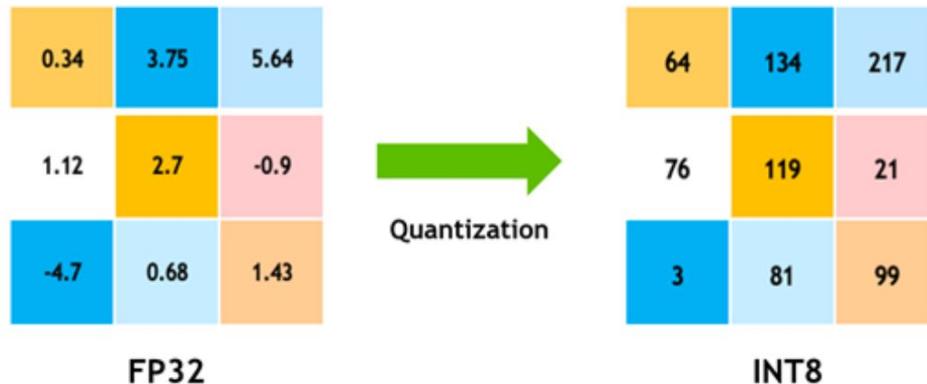
Risks from Synthetic Data

Model collapse setting



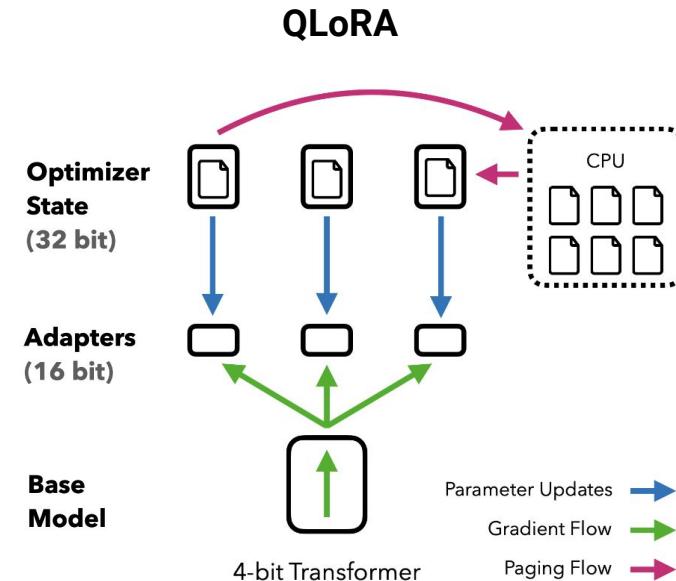
I. Shumailov et al. [AI Models Collapse When Trained on Recursively Generated Data](#). 2024.

Quantization



Model parameters can be stored in fewer bits. (FP32→INT8)

Even fewer: 4 bits



Interpretability / Representation Engineering

Most of the 4,096 features found by the auto-encoder have consistent, interpretable responses to input data

#3647 - Legal language


```
and to deposit the same in Marshall's checking  
upon the intention, actual or presumed, of  
mortgage which was then existing on the property,  
be permitted to retain actual possession of the mo  
lien on encumbered property of the estate
```

#2937 - DNA Sequences

```
AGTTTCGTTTACATG GGG  
AGACAACCTTTCTTT Ex3  
ACACACGGACACGGCTACGG  
CTCCGTGTTTGMDM2-  
CAAGAAAAACCATGCTTGTT
```

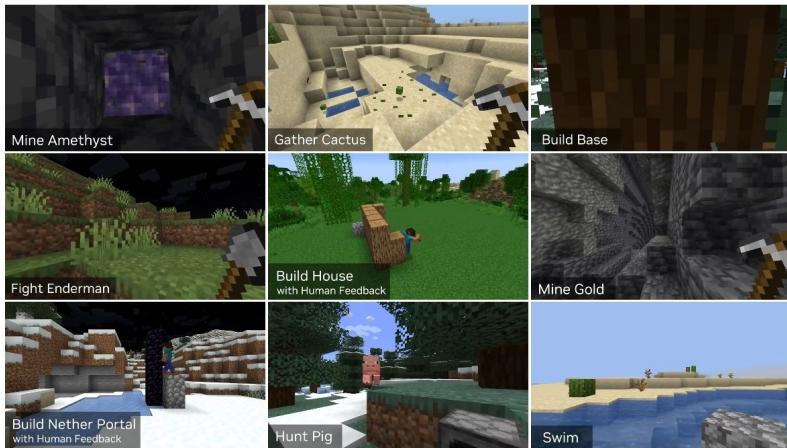
#28 - Code: HTTP requests and responses

```
Net :: HttpResponseMessage #{$ response . code  
HTTPMethod { case "GET": id :=  
405, HttpResponseMessage500, ### HTTP  
setHTTPMethod: @"POST"];@{  
atoKey, forHTTPHeaderField: "
```



A visualization of a transformer layer with 512 neurons decomposed into more than 4000 features with semantic meaning.

Beyond Text: Robotics, Simulations, Physical Tasks




LLM-powered embodied lifelong learning agent in Minecraft that continuously explores the world, acquires diverse skills, and makes novel discoveries without human intervention.

G. Wang et al. [Voyager: An Open-Ended Embodied Agent with Large Language Models](#). 2023.

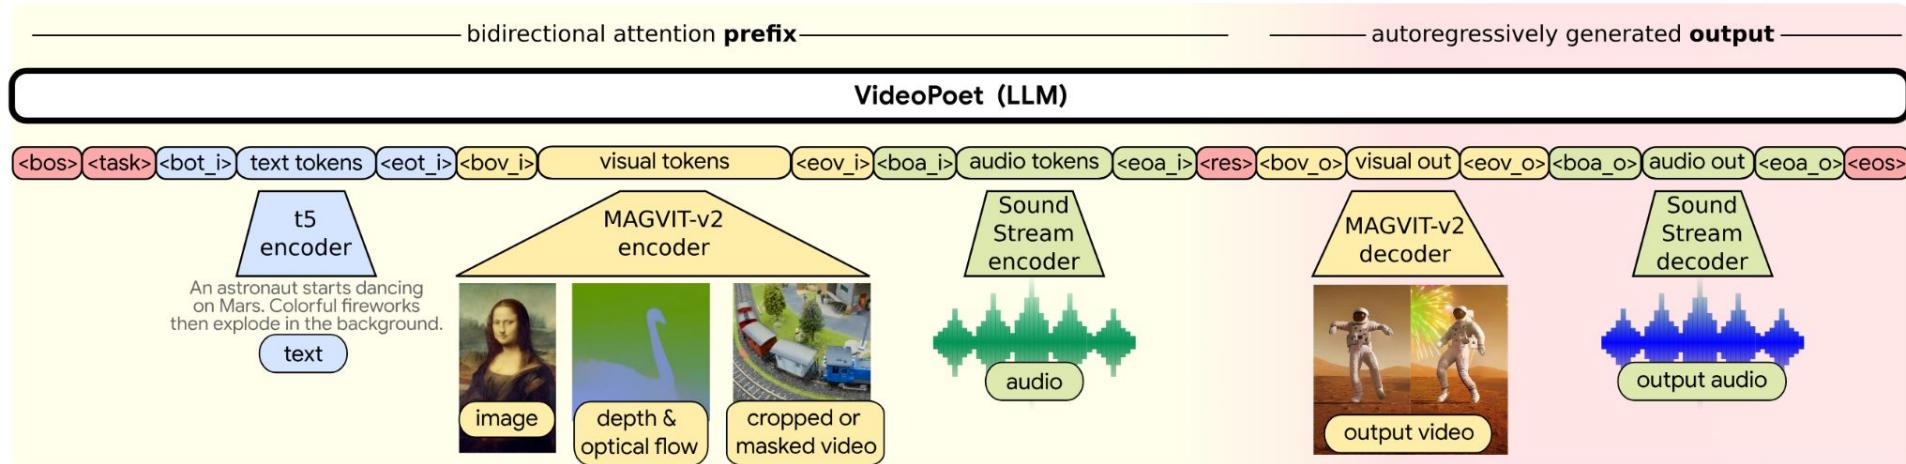
Mobile ALOHA doing complex long-horizon tasks. It extends ALOHA with a mobile base and a whole-body teleoperation interface. ALOHA learns a policy using transformers to predict a sequence of actions.

Z. Fu et al. [Mobile ALOHA: Learning Bimanual Mobile Manipulation with Low-Cost Whole-Body Teleoperation](#). 2024.



Beyond Text: Video Generation

Autoregressively generate tokens of multiple modalities (images, videos, text, audio).



K. Dan et al. [VideoPoet: A Large Language Model for Zero-Shot Video Generation](#). 2024.

Retrieval-Augmented Generation (RAG)

LLMs have no knowledge beyond training date, and frequent updates to model are impractical.

Idea: Augment LLMs with *retrieval system*!

Example: Retrieval Step of RAG

**Struggle for Rome
(board game)** —
Catan Histories:
Struggle for Rome is a
2006 German-style
board game based on
the game mechanics of
"Settlers of Catan",
depicting the fall of the
Roman Empire...

The Kids of Catan
— The Kids of Catan
is a German board
game designed for
children using the
theme from "The
Settlers of Catan"...

Pirate's Cove —
Pirate's Cove is a
board game designed
by Paul Randles and
Daniel Stahl, originally
published in Germany
in 2002...

**Catan: Cities &
Knights** — Catan:
Cities & Knights,
formerly "The Cities
and Knights of Catan"
is an expansion to the
board game "The
Settlers of Catan"...

Catan — The Settlers
of Catan, sometimes
shortened to Catan or
Settlers, is a
multiplayer board
game designed by
Klaus Teuber and first
published in 1995...

Question:

Which board game was published most
recently, Pirate's Cove or Catan?



Vector DB

Example: Generation Step of RAG

Question:

Which board game was published most recently, Pirate's Cove or Catan?

Pirate's Cove —

Pirate's Cove is a board game designed by Paul Randles and Daniel Stahl, originally published in Germany in 2002...

Catan —

The Settlers of Catan, sometimes shortened to Catan or Settlers, is a multiplayer board game designed by Klaus Teuber and first published in 1995...




LLM



Answer:
Pirate's Cove



Example: Without RAG



Answer:
The board game that was published most recently is [Settlers of] Catan.


Retrieval-Augmented Generation (RAG)

LLMs have no knowledge beyond training date, and frequent updates to model are impractical.

Idea: Augment LLMs with *retrieval system*!

But at scale, sensitive to choices of:

- 1) chunking strategy,
- 2) embedding model, and
- 3) generation model.




And not guaranteed to be hallucination-free ...

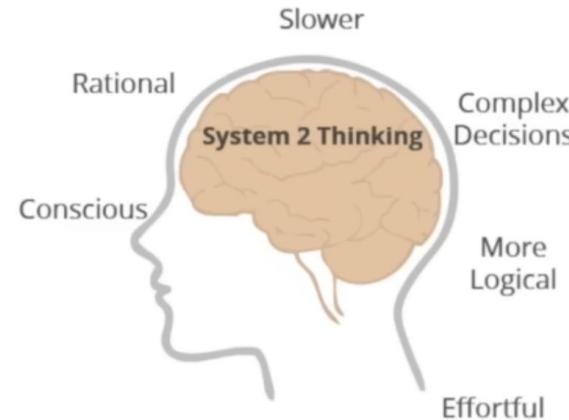
Two Modes of Human Thought

Prompting and Fine-Tuning can still only yield a (better) System 1

System 1

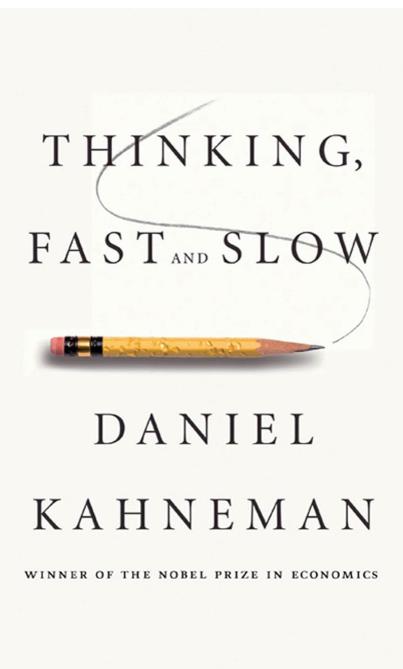


System 2



$$2 + 2 =$$

$$17 \times 24 =$$



Neurosymbolic To Combine Both Worlds

Deep Learning [System 1]

- Sub-symbolic knowledge
- Open-domain knowledge
- Rapid reasoning
- Handling noise and naturalness
- In-context learning

Classical Algorithms [System 2]

- Domain-specific knowledge
- Complex reasoning
- Interpretability
- Compositional reasoning
- Generalizability

neural



symbolic



neurosymbolic

Example: Extracting Knowledge Using GPT

mountain

name
Everest
Fuji
K2
Mt. Blanc

```
@gpt("The height of {{x}} is {{y}} in meters")
type height(bound x: String, free y: i32)
// Retrieving height of mountains
rel mount_height(m, h) = mountain(m) and height(m, h)
```

mountain_height

name	height
Everest	8848
Fuji	3776
K2	8611
Mt. Blanc	4808

mountain names come from
a database, which cannot
hallucinate!

Example: Classifying Images Using CLIP

image	
id	image
0	
1	
...	...

```
@clip(["cat", "dog"])
type classify(bound img: Tensor, free label: String)
// Classify each image as cat or dog
rel cat_or_dog(i, l) = image(i, m) and classify(m, l)
```

cat_or_dog		
prob	id	label
0.00	0	cat
0.99	0	dog
0.98	1	cat
0.02	1	dog
...

What Is The Future: Key Topics

- Synthetic Data / Distillation
- Model Compression
- Interpretability
- Beyond Text: Robotics, Video, etc.
- RAG and Vector Databases
- Agent Frameworks
- Neurosymbolic Learning

Course Overview and Pre-requisites

In-depth exploration of large language models (LLMs) with a focus on designing, training, and using them.

- Start with design decisions behind attention mechanism and transformer architectures.
- progress through practical aspects of pre-training and efficient deployment at scale.
- culminate in usage techniques such as prompting, RAG, and neuro-symbolic learning.

Pre-requisites:

1. CIS 5200 or equivalent: Mathematical foundations of machine learning.
2. CIS 5450 or equivalent: Experience with building, training, and debugging machine learning models.

Scope of Course

Areas of emphasis:

- **Foundations:** lectures cover broadly applicable and (relatively) established techniques
- **Systems:** homeworks implementing those techniques using deep learning frameworks
- **Research:** topics derived from recent papers in top ML conferences (NeurIPS/ICLR/ICML)
- **Experimentation:** team project to implement and empirically evaluate a new technique

Topics not covered:

- **Application Domains:** we won't dive into specific domains like NLP, Vision, or Robotics
- **Theory:** limited to mathematical concepts needed to understand and implement techniques
- **Classical ML:** we won't cover classical ML approaches that predate LLMs
- **AI Application Dev:** we won't teach you the AI dev stack or how to build enterprise AI apps

Learning Objectives

- Analyze design decisions in modern and upcoming transformer architectures.
- Determine the hardware, software, and data requirements for pre-training or fine-tuning an LLM for new tasks.
- Understand where LLMs should and should not be used based on their capability and reliability.
- Leverage a deep understanding of LLM theory and software to design prompts and applications around them.

Course Activities

- Lectures by instructor, external guest speakers (virtual and in-person), and local guest speakers (TAs or PhD students). Lectures will not be recorded!
- Five challenging programming assignments to be done individually (**55%**).
- Project (**25%**): Deep-dive into implementing and analyzing an LLM technique in teams of 2-3 students.
- Final Exam (**15%**): Any concepts covered in the lectures.
- Class Participation (**5%**): You are expected to attend all lectures and especially guest lectures; we will measure class participation by taking attendance on randomly chosen days.

Homeworks

- HW0:** Introductory assignment comparing and analyzing outputs from different LLMs.
- HW1:** Build and understand the Transformer architecture from the ground up.
- HW2:** Explore techniques to adapt pre-trained LLMs to new tasks in an efficient and performant manner.
- HW3:** Leverage patterns in pretrained weights to compress LLMs for memory-efficient inference and fine-tuning.
- HW4:** Investigate the intersection of LLMs with symbolic reasoning and apply it to challenging reasoning tasks.

Course Resources

Official Website: <https://llm-class.github.io/>

Also: Canvas, Gradescope, Ed Discussions.